

PART 2B. SONG POPULARITY

GROUP 9

Importing Libraries

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
import plotly.express as px

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC, LinearSVC
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

Loading the dataset

```
In [ ]: song_path = r'spotifydata.csv'
song_data = pd.read_csv(song_path)
```

```
In [ ]: print("EDA:")
print('\nSong Data:')
display(song_data.head())
print("Length of song dataset:", len(song_data))
display(song_data.describe())
display(song_data.isna().sum())
display(pd.DataFrame({'Column': song_data.columns, 'DType': [song_data[col].dtype
```

EDA:

Song Data:

| | genre | artist_name | track_name | track_id | popularity | acousticness | danceabi |
|---|-------|-------------------|----------------------------------|------------------------|------------|--------------|----------|
| 0 | Movie | Henri Salvador | C'est beau de faire un Show | 0BRjO6ga9RKCKjfDqeFgWV | 0 | 0.611 | 0. |
| 1 | Movie | Martin & les fées | Perdu d'avance (par Gad Elmaleh) | 0BjC1NfoEOOusryehmNudP | 1 | 0.246 | 0. |
| 2 | Movie | Joseph Williams | Don't Let Me Be Lonely Tonight | 0CoSDzoNIKCRs124s9uTVy | 3 | 0.952 | 0. |
| 3 | Movie | Henri Salvador | Dis-moi Monsieur Gordon Cooper | 0Gc6TVm52BwZD07Ki6tlvf | 0 | 0.703 | 0. |
| 4 | Movie | Fabien Nataf | Ouverture | 0lusIXpMROHdEPvSI1ftQK | 4 | 0.950 | 0. |

Length of song dataset: 232725

| | popularity | acousticness | danceability | duration_ms | energy | instrumentaln |
|-------|---------------|---------------|---------------|--------------|---------------|---------------|
| count | 232725.000000 | 232725.000000 | 232725.000000 | 2.327250e+05 | 232725.000000 | 232725.0000 |
| mean | 41.127502 | 0.368560 | 0.554364 | 2.351223e+05 | 0.570958 | 0.1483 |
| std | 18.189948 | 0.354768 | 0.185608 | 1.189359e+05 | 0.263456 | 0.3027 |
| min | 0.000000 | 0.000000 | 0.056900 | 1.538700e+04 | 0.000020 | 0.0000 |
| 25% | 29.000000 | 0.037600 | 0.435000 | 1.828570e+05 | 0.385000 | 0.0000 |
| 50% | 43.000000 | 0.232000 | 0.571000 | 2.204270e+05 | 0.605000 | 0.0000 |
| 75% | 55.000000 | 0.722000 | 0.692000 | 2.657680e+05 | 0.787000 | 0.0358 |
| max | 100.000000 | 0.996000 | 0.989000 | 5.552917e+06 | 0.999000 | 0.9990 |

genre0

artist_name0

track_name0

track_id0

popularity0

acousticness0

danceability0

duration_ms0

energy0

instrumentalness0

key0

liveness0

loudness0

mode0

speechiness0

tempo0

time_signature0

valence0

dtype: int64

| | Column | DType | NUniques |
|----|------------------|---------|----------|
| 0 | genre | object | 27 |
| 1 | artist_name | object | 14564 |
| 2 | track_name | object | 148615 |
| 3 | track_id | object | 176774 |
| 4 | popularity | int64 | 101 |
| 5 | acousticness | float64 | 4734 |
| 6 | danceability | float64 | 1295 |
| 7 | duration_ms | int64 | 70749 |
| 8 | energy | float64 | 2517 |
| 9 | instrumentalness | float64 | 5400 |
| 10 | key | object | 12 |
| 11 | liveness | float64 | 1732 |
| 12 | loudness | float64 | 27923 |
| 13 | mode | object | 2 |
| 14 | speechiness | float64 | 1641 |
| 15 | tempo | float64 | 78512 |
| 16 | time_signature | object | 5 |
| 17 | valence | float64 | 1692 |

What did we do in the above step? - The code is about Exploratory Data Analysis (EDA) on the dataset. It displays first few data points, its length, statistical summary, count of missing values per column, and a tabulation displaying column names, their data types, and the number of unique values within each column.

```
In [ ]: # Retrieve the keys present in the song_data dictionary
song_data.keys()
```

```
Out[ ]: Index(['genre', 'artist_name', 'track_name', 'track_id', 'popularity',
          'acousticness', 'danceability', 'duration_ms', 'energy',
          'instrumentalness', 'key', 'liveness', 'loudness', 'mode',
          'speechiness', 'tempo', 'time_signature', 'valence'],
          dtype='object')
```

Checking Missing Values

```
In [ ]: song_data.isna().sum()
```

```
Out[ ]: genre          0
artist_name         0
track_name          0
track_id            0
popularity           0
acousticness        0
danceability         0
duration_ms         0
energy              0
instrumentalness     0
key                 0
liveness            0
loudness            0
mode                0
speechiness         0
tempo              0
time_signature      0
valence             0
dtype: int64
```

Data Visualization

```
In [ ]: sns.distplot(song_data['popularity']).set_title('Popularity Distribution')
```


C:\Users\patel\AppData\Local\Temp\ipykernel_12888\3871686646.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
Out[ ]: sns.distplot(song_data['popularity']).set_title('Popularity Distribution')
Text(0.5, 1.0, 'Popularity Distribution')
```




What did we do in the above step? - The code is to plot the distribution plot (distplot) for the 'popularity' column. Using this, we are visualizing the distribution of popularity scores among the songs.

Visualizing Correlation Among Song Data Features Using Heatmap

```
In [ ]: correlation = song_data.corr()
fig = plt.figure(figsize=(20,5))
mask = np.triu(np.ones_like(correlation, dtype=bool), k=1)
sns.heatmap(correlation, annot=True, cmap="coolwarm", mask=mask)
plt.show()
```

C:\Users\patel\AppData\Local\Temp\ipykernel_12888\2112152709.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

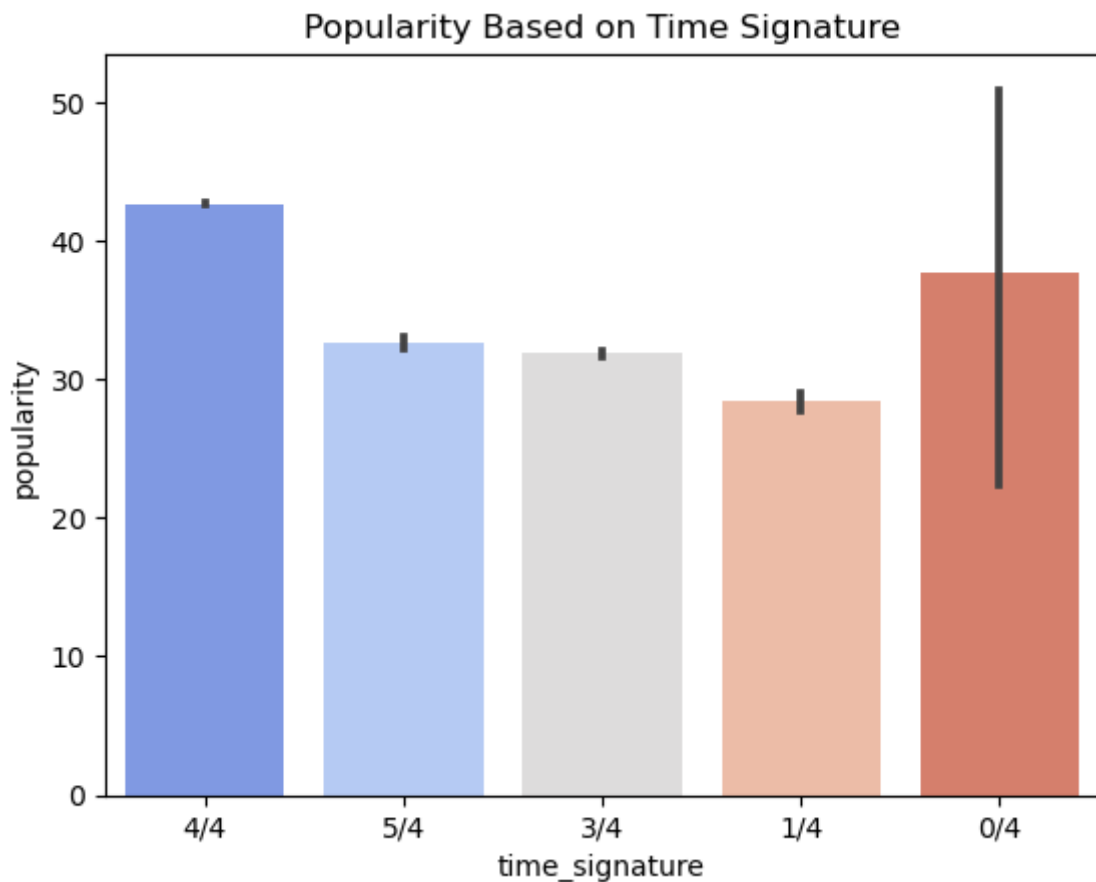
```
correlation = song_data.corr()
```



Visualization of Popularity Based on Time Signature

```
In [ ]: sns.barplot(x = 'time_signature', y = 'popularity', data = song_data, palette='coolwarm')
plt.title('Popularity Based on Time Signature')
```

```
Out[ ]: Text(0.5, 1.0, 'Popularity Based on Time Signature')
```

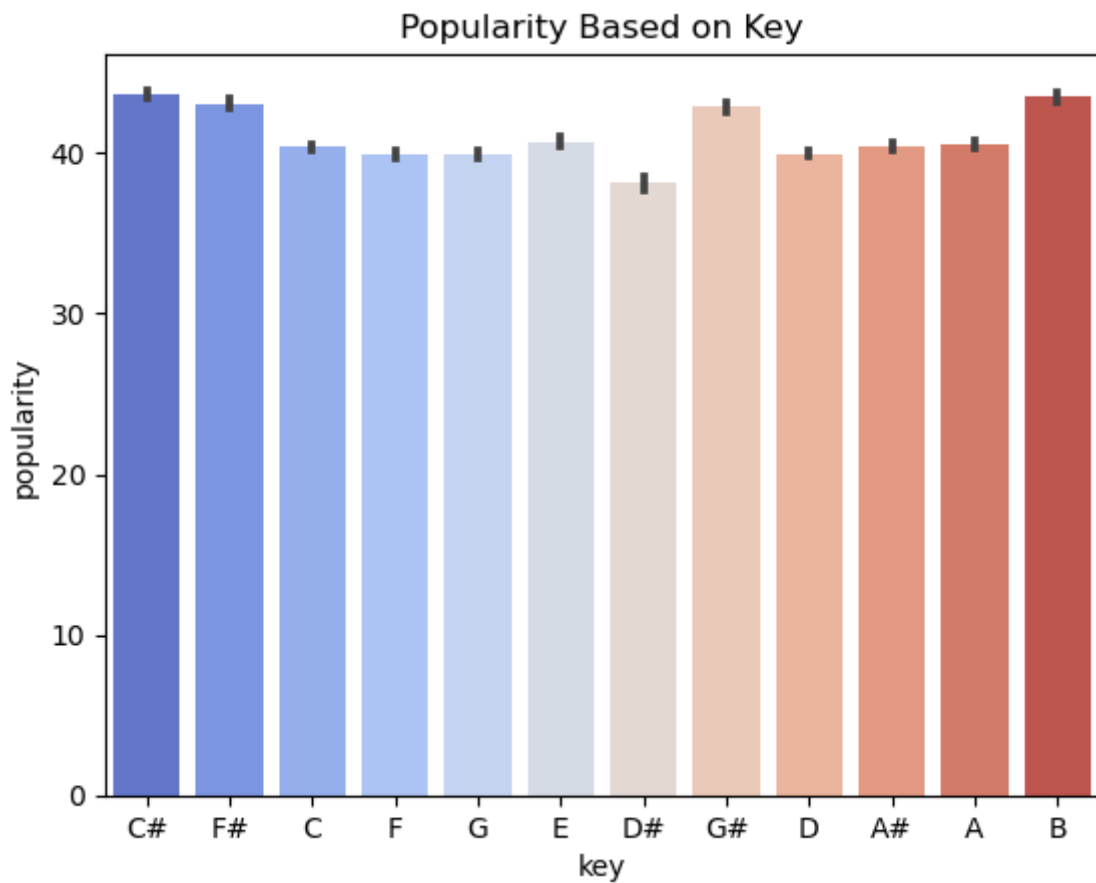


What did we do in the above step? - We created a bar plot to showcase the relationship between the time signature of songs and their respective popularity levels

Popularity Distribution Across Different Musical Keys

```
In [ ]: sns.barplot(x = 'key', y = 'popularity', data = song_data, palette='coolwarm')
plt.title('Popularity Based on Key')
```

```
Out[ ]: Text(0.5, 1.0, 'Popularity Based on Key')
```




What did we do in the above step? - We created a bar plot to show how popularity varies across different musical keys, mainly to understand relationships between key and music popularity.

Popularity Based on Musical Mode

```
In [ ]: sns.barplot(x = 'mode', y = 'popularity', data = song_data, palette='coolwarm')  
plt.title('Popularity Based on Mode')
```

```
Out[ ]: Text(0.5, 1.0, 'Popularity Based on Mode')
```



What did we do in the above step? - We created a bar plot to show the relationship between musical mode and song popularity

Distribution of Acousticness in Songs with Popularity Above 70

```
In [ ]: popular_songs = song_data[song_data.popularity > 70]
sns.distplot(popular_songs['acousticness'])
plt.title('Acousticness for Songs with More than 70 Popularity')
```

C:\Users\patel\AppData\Local\Temp\ipykernel_12888\1018885752.py:2: UserWarning:

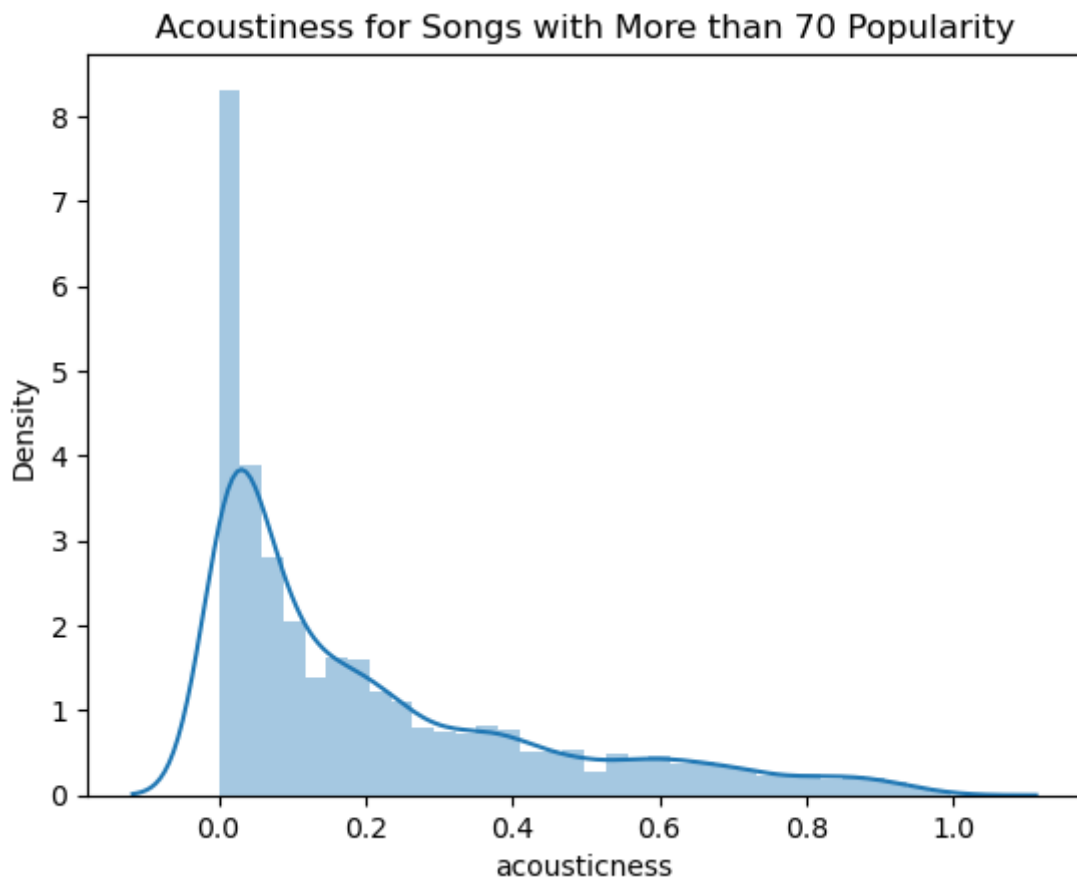
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(popular_songs['acousticness'])
```

```
Out[ ]: Text(0.5, 1.0, 'Acousticness for Songs with More than 70 Popularity')
```

What did we do in the above step? - The code filters songs with a popularity score above 70 from the dataset and visualization is performed specifically for acoustiness

Distribution of Acoustiness in Songs with Popularity Below 70

```
In [ ]: unpopular_songs = song_data[song_data.popularity < 70]
sns.distplot(unpopular_songs['acoustiness'])
plt.title('Acoustiness for Songs with Less than 50 Popularity')
```

C:\Users\patel\AppData\Local\Temp\ipykernel_12888\228728011.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(unpopular_songs['acoustiness'])
```

```
Out[ ]: Text(0.5, 1.0, 'Acoustiness for Songs with Less than 50 Popularity')
```



What did we do in the above step? - The code filters songs with a popularity score below 70 from the dataset and visualization is performed specifically for acoustiness

Distribution of Loudness for Highly Popular Songs

```
In [ ]: sns.distplot(popular_songs['loudness'])
plt.title('Loudness for Songs with More than 70 Popularity')
```

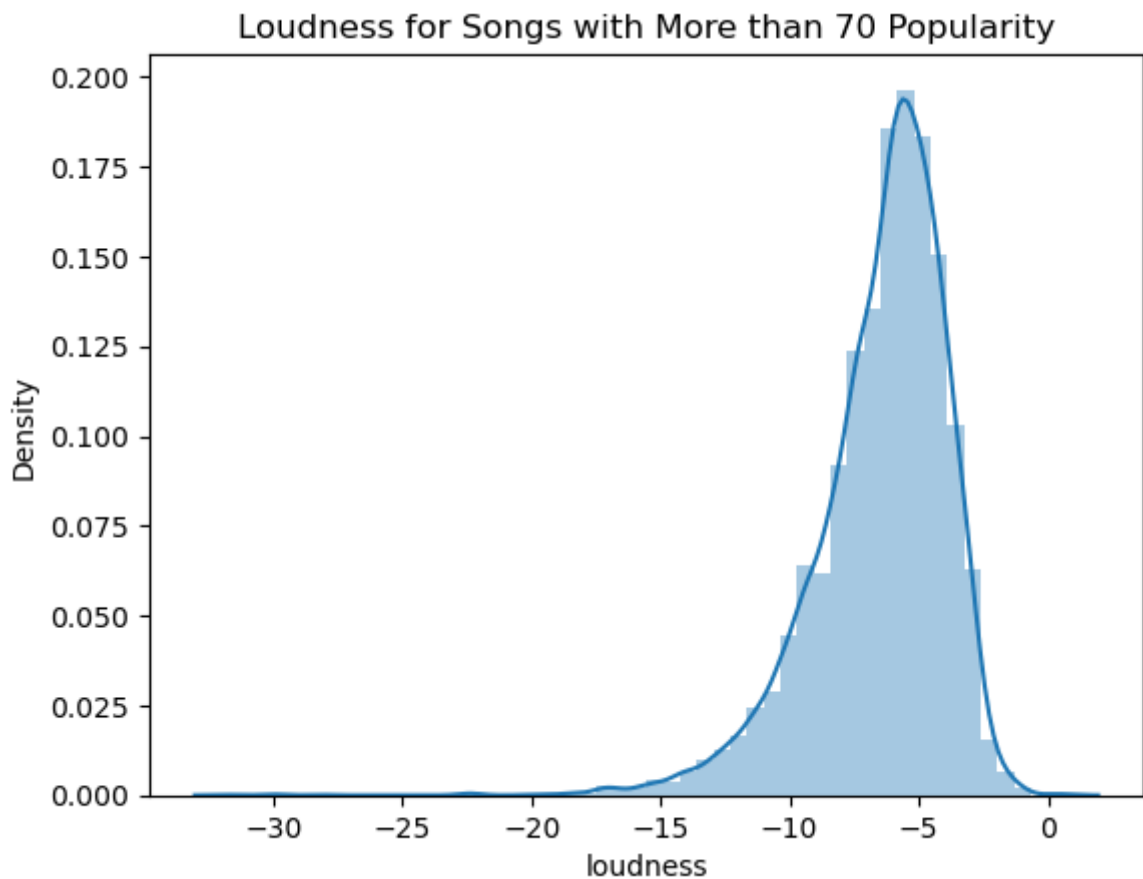
C:\Users\patel\AppData\Local\Temp\ipykernel_12888\1995242145.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(popular_songs['loudness'])
Out[ ]: Text(0.5, 1.0, 'Loudness for Songs with More than 70 Popularity')
```



What did we do in the above step? - Visualizing the distribution of loudness for songs with a popularity above 70, giving information about the loudness patterns within the list of highly popular songs.

Distribution of Loudness for Less Popular Songs

```
In [ ]: unpopular_songs = song_data[song_data.popularity < 70]
sns.distplot(unpopular_songs['loudness'])
plt.title('Loudness for Songs with Less than 70 Popularity')
```

C:\Users\patel\AppData\Local\Temp\ipykernel_12888\1000170378.py:2: UserWarning:

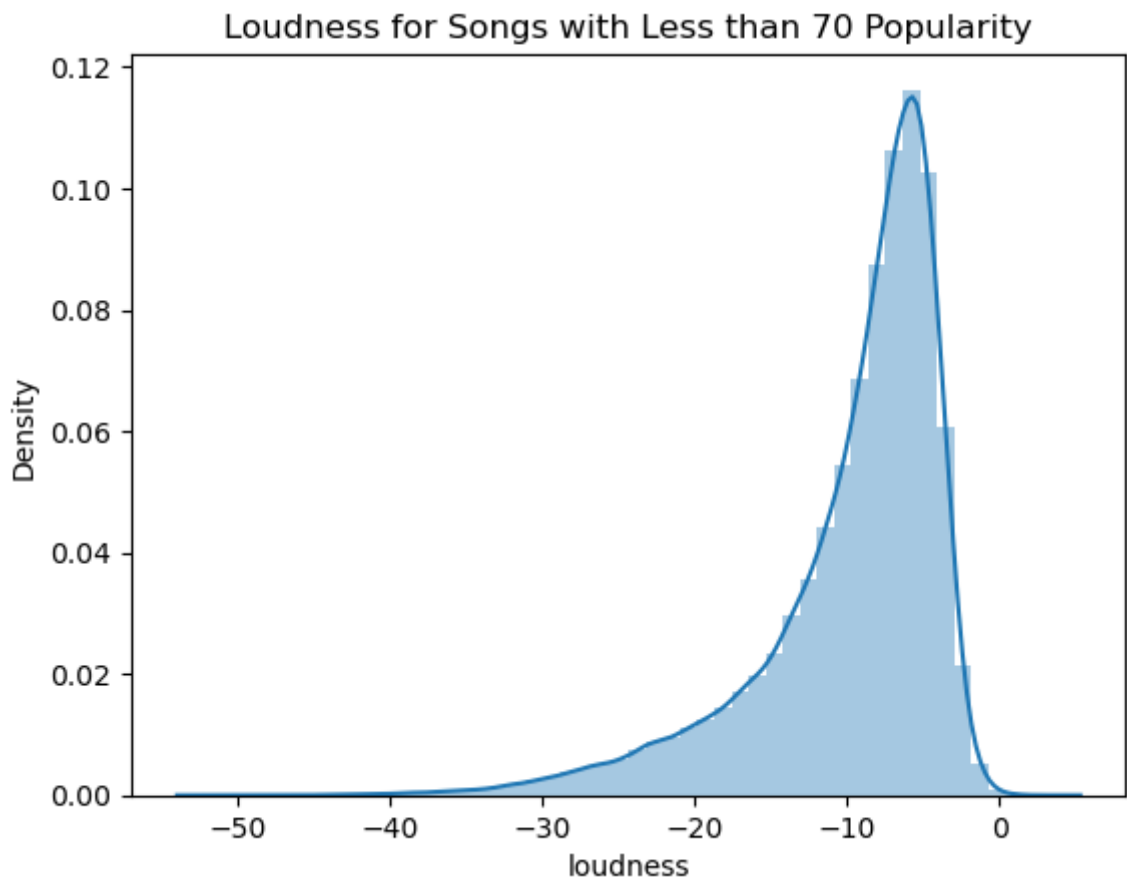
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(unpopular_songs['loudness'])
```

```
Out[ ]: Text(0.5, 1.0, 'Loudness for Songs with Less than 70 Popularity')
```



What did we do in the above step? - Visualizing the distribution of loudness for songs with a popularity below 70, giving information about the loudness patterns within the list of less popular songs.

Data Preparation

Encoding Categorical Keys in Song Data

```
In [ ]: unique_keys = song_data['key'].unique()
for i in range(len(unique_keys)):
    song_data.loc[song_data['key'] == unique_keys[i], 'key'] = i
song_data.sample(5)
```

| | genre | artist_name | track_name | track_id | popularity | acousticness |
|---------------|-------------|--------------------|------------------------|------------------------|------------|--------------|
| 48743 | Blues | Joe Cocker | Woman To Woman | 0K3Dem9aRwRG9DVo0W2T3o | 34 | 0.178 |
| 6842 | Alternative | Ayelle | Obvious | 2ySyV2K8WnOqsQ0CfBNdnx | 54 | 0.284 |
| 68302 | Hip-Hop | Tyler, The Creator | I Ain't Got Time! | 430qNtapCS3Ue1yoSql1oV | 65 | 0.053 |
| 109276 | Pop | The Foundations | Build Me Up Buttercup | 6sPOmDulFtLzfX25zICNrC | 74 | 0.313 |
| 210078 | World | Rend Collective | Weep With Me - Reprise | 4sYJN5wnWsiuFfYPCbuuF4 | 37 | 0.782 |

Mapping Time Signatures to Numerical Categories in Song Data

```
In [ ]: time_sign = song_data['time_signature'].unique()
for i in range(len(time_sign)):
    song_data.loc[song_data['time_signature'] == time_sign[i], 'time_signature']
song_data.sample(5)
```

| | genre | artist_name | track_name | track_id | popularity | acousticness |
|---------------|---------|---------------|--------------------------------|------------------------|------------|--------------|
| 8137 | Country | Jerry Garcia | Sugaree | 4XoYeoIVYTiddO9wZLXLgl | 44 | 0.0914 |
| 176845 | Jazz | Johnny Mathis | The Last Time I Felt Like This | 2hZt9BmpFVRZq5xzOeKSA7 | 45 | 0.7730 |
| 7876 | Country | Roger Creager | Love | 7zZKjQDm8JNFkrxRvafHid | 36 | 0.1920 |
| 61438 | R&B | Sex on Toast | Oh, Loretta! | 4HqxTpbXpLO1jgKJppk8GI | 52 | 0.3470 |
| 7105 | Country | Alan Jackson | Like Red On a Rose | 1ayFArNqsYgGT8gWWScTD | 37 | 0.6340 |

Converting Musical Modes to Numerical Representation

```
In [ ]: song_data.loc[song_data["mode"] == 'Major', "mode"] = 1
song_data.loc[song_data["mode"] == 'Minor', "mode"] = 0
song_data.sample(5)
```

Out[]:

| | genre | artist_name | track_name | track_id | popularity | acousticness |
|---------------|-----------|--------------------|-----------------------------|-------------------------|------------|--------------|
| 228826 | Soul | Teena Marie | I Need Your Lovin' | 0HYCfsLMVE61LZIVI3d1UA | 44 | 0.00889 |
| 32360 | Anime | SEKAI NO OWARI | Rafflesia | 67abPmmiZMrgEW7UFnjJfF | 19 | 0.05080 |
| 217548 | World | Hildur Guðnadóttir | Elevation | 3x4AbXBLj5t3x1OnAAVBiL | 37 | 0.86100 |
| 158585 | Reggaeton | Brytiago | Dime a Vel (feat. Almighty) | 65fV6PGE2bytfHrpPgZPPb | 49 | 0.30300 |
| 54919 | R&B | The Weeknd | Reminder - Remix | 7F1LWA9sII TorHUo4amGqk | 64 | 0.14900 |

What did we do in the above step? - Transforming categorical labels 'Major' and 'Minor' to numerical representations '1' and '0' respectively.

Binary Classification of Song Popularity

```
In [ ]: song_data.loc[song_data['popularity'] < 70, 'popularity'] = 0
song_data.loc[song_data['popularity'] >= 70, 'popularity'] = 1
song_data.loc[song_data['popularity'] == 1]
```

| Out[]: | genre | artist_name | track_name | track_id | popularity | acousticness | c |
|---------|-------|------------------|---|------------------------|------------|--------------|-----|
| 145 | R&B | Mariah Carey | Hero | 4FCb4CUBFCMNRki6IYc1zI | 1 | 0.7350 | |
| 147 | R&B | Jason Derulo | Tip Toe (feat. French Montana) | 2z4pcBLQXF2BXKFvd0BuB6 | 1 | 0.0233 | |
| 160 | R&B | Jennifer Lopez | Dinero | 22mQXNE0nCuWq4yOwcadln | 1 | 0.4100 | |
| 161 | R&B | Rihanna | Hate That I Love You | 7iu0WYLD04yKsKf3seaxzI | 1 | 0.3230 | |
| 174 | R&B | Usher | OMG (feat. will.i.am) | 1bM50INir8voAkVoKuvEUI | 1 | 0.1980 | |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 226102 | Rock | Bacilos | Caraluna | 4XTMj7kd8DHLI0r7ghmEAr | 1 | 0.1760 | |
| 226281 | Rock | Alejandro Sanz | Corazón partío | 0wQCKR9OFjYu5Kzrk7WivJ | 1 | 0.1990 | |
| 226307 | Rock | Elefante | Así Es La Vida | 3ge3q3Hz0KWhQX5EAQcwEy | 1 | 0.1310 | |
| 226413 | Rock | Roxette | It Must Have Been Love - From the "Pretty Woma... | 6qB7YcFpeBEQa0D6QO482y | 1 | 0.3400 | |
| 226424 | Rock | La Mosca Tse-Tse | Para No Verte Más | 19CmuECYssqkPWANF4nLWM | 1 | 0.0168 | |

9001 rows × 18 columns

What did we do in the above step? - Transforming 'popularity' column into a binary classification by setting unpopular songs (below 70) to 0 and popular songs (above 70) to 1. The final line includes only songs marked as highly popular (assigned a value of 1).
- It is our threshold that we consider only songs above 70 as popular songs and those below 70 as unpopular songs

Model Building

Feature Selection

```
In [ ]: song_features = ["acousticness", "danceability", "duration_ms", "energy", "instru",
                        "mode", "speechiness", "tempo", "time_signature", "valence"]
```

What did we do in the above step? - For feature selection, we will select the following features which are only based on music theory and not artist/song information:

acousticness, danceability, duration_ms, energy, instrumentalness, key, liveliness, loudness, mode, speechiness, tempo, time_signature, and valence.

```
In [ ]: training = song_data.sample(frac = 0.8, random_state = 420)
X_train = training[song_features]
y_train = training['popularity']
X_test = song_data.drop(training.index)[song_features]
```

```
In [ ]: X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train, test_size
```

```
In [ ]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_valid = scaler.transform(X_valid)
```

Model 1: Logistic Regression

```
In [ ]: lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)
lr_y_pred = lr_model.predict(X_valid)
lr_accuracy = accuracy_score(y_valid, lr_y_pred)
print("Accuracy: " + str(lr_accuracy))

lr_auc = roc_auc_score(y_valid, lr_y_pred)
print("AUC: " + str(lr_auc))
```

Accuracy: 0.9600655279836717

AUC: 0.5

```
In [ ]: from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, re

accuracy = accuracy_score(y_valid, lr_y_pred)
conf_matrix = confusion_matrix(y_valid, lr_y_pred)
precision = precision_score(y_valid, lr_y_pred, average='weighted')
recall = recall_score(y_valid, lr_y_pred, average='weighted')
f1 = f1_score(y_valid, lr_y_pred, average='weighted')

tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)

print("Confusion Matrix:")
print(conf_matrix)
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F-measure: {f1:.2f}")
print(f"Specificity: {specificity:.2f}")

plt.figure(figsize=(4, 3))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Low', 'Predicted High'],
            yticklabels=['Actual Low', 'Actual High'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Logistic Regression')
plt.show()

lr_y_pred_proba = lr_model.predict_proba(X_valid)[: , 1]

fpr, tpr, thresholds = roc_curve(y_valid, lr_y_pred_proba)
```



```

roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - Logistic Regression')
plt.legend(loc='lower right')
plt.show()

```

c:\Users\patel\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:146
 9: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in label
 s with no predicted samples. Use `zero_division` parameter to control this behavi
 or.

```
_warn_prf(average, modifier, msg_start, len(result))
```

Confusion Matrix:

```
[[35749    0]
 [ 1487    0]]
```


Accuracy: 0.96


Precision: 0.92

Recall: 0.96

F-measure: 0.94

Specificity: 1.00





Model 2: Decision Tree Classifier

```
In [ ]: dt_model = DecisionTreeClassifier()
dt_model.fit(X_train, y_train)
dt_y_pred = dt_model.predict(X_valid)
dt_accuracy = accuracy_score(y_valid, dt_y_pred)
print("Accuracy: " + str(dt_accuracy))

dt_auc = roc_auc_score(y_valid, dt_y_pred)
print("AUC: " + str(dt_auc))
```

Accuracy: 0.9671822967021162
AUC: 0.8430473071768053

```
In [ ]: from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, re
import seaborn as sns
import matplotlib.pyplot as plt

accuracy = accuracy_score(y_valid, dt_y_pred)
conf_matrix = confusion_matrix(y_valid, dt_y_pred)
precision = precision_score(y_valid, dt_y_pred, average='weighted')
recall = recall_score(y_valid, dt_y_pred, average='weighted')
f1 = f1_score(y_valid, dt_y_pred, average='weighted')

tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)

print("Confusion Matrix:")
print(conf_matrix)
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
```

```

print(f"F-measure: {f1:.2f}")
print(f"Specificity: {specificity:.2f}")

plt.figure(figsize=(4, 3))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Low', 'Predicted High'],
            yticklabels=['Actual Low', 'Actual High'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Decision Tree')
plt.show()

dt_y_pred_proba = dt_model.predict_proba(X_valid)[: , 1]

fpr, tpr, thresholds = roc_curve(y_valid, dt_y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - Decision Tree')
plt.legend(loc='lower right')
plt.show()

```

Confusion Matrix:

```
[[34961  788]
 [ 434 1053]]
```


Accuracy: 0.97


Precision: 0.97

Recall: 0.97

F-measure: 0.97

Specificity: 0.98





Model 3: Random Forest Classifier

```
In [ ]: rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
rf_y_pred = rf_model.predict(X_valid)
rf_accuracy = accuracy_score(y_valid, rf_y_pred)
print("Accuracy: " + str(rf_accuracy))

rf_auc = roc_auc_score(y_valid, rf_y_pred)
print("AUC: " + str(rf_auc))
```

Accuracy: 0.986196154259319
AUC: 0.8429595888075875

```
In [ ]: from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score

accuracy = accuracy_score(y_valid, rf_y_pred)
conf_matrix = confusion_matrix(y_valid, rf_y_pred)
precision = precision_score(y_valid, rf_y_pred, average='weighted')
recall = recall_score(y_valid, rf_y_pred, average='weighted')
f1 = f1_score(y_valid, rf_y_pred, average='weighted')

tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)

print("Confusion Matrix:")
print(conf_matrix)
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F-measure: {f1:.2f}")
print(f"Specificity: {specificity:.2f}")
```

```

plt.figure(figsize=(4, 3))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Low', 'Predicted High'],
            yticklabels=['Actual Low', 'Actual High'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Random Forest')
plt.show()

rf_y_pred_proba = rf_model.predict_proba(X_valid)[: , 1]

fpr, tpr, thresholds = roc_curve(y_valid, rf_y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - Random Forest')
plt.legend(loc='lower right')
plt.show()

```

Confusion Matrix:

```
[[35700   49]
 [  465 1022]]
```


Accuracy: 0.99


Precision: 0.99

Recall: 0.99

F-measure: 0.99

Specificity: 1.00





Model 4: K-Nearest Neighbors Classifier

```
In [ ]: knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
knn_y_pred = knn_model.predict(X_valid)
knn_accuracy = accuracy_score(y_valid, knn_y_pred)
print("Accuracy: " + str(knn_accuracy))

knn_auc = roc_auc_score(y_valid, knn_y_pred)
print("AUC: " + str(knn_auc))
```

Accuracy: 0.9545063916639811
AUC: 0.5699358166028053

```
In [ ]: from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, re
import matplotlib.pyplot as plt
import seaborn as sns

accuracy = accuracy_score(y_valid, knn_y_pred)
conf_matrix = confusion_matrix(y_valid, knn_y_pred)
precision = precision_score(y_valid, knn_y_pred, average='weighted')
recall = recall_score(y_valid, knn_y_pred, average='weighted')
f1 = f1_score(y_valid, knn_y_pred, average='weighted')

tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)

print("Confusion Matrix:")
print(conf_matrix)
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
```

```

print(f"F-measure: {f1:.2f}")
print(f"Specificity: {specificity:.2f}")

plt.figure(figsize=(4, 3))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Low', 'Predicted High'],
            yticklabels=['Actual Low', 'Actual High'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - KNN')
plt.show()

knn_y_pred_proba = knn_model.predict_proba(X_valid)[: , 1]

fpr, tpr, thresholds = roc_curve(y_valid, knn_y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - KNN')
plt.legend(loc='lower right')
plt.show()

```

Confusion Matrix:

```
[[35316  433]
 [ 1261   226]]
```


Accuracy: 0.95

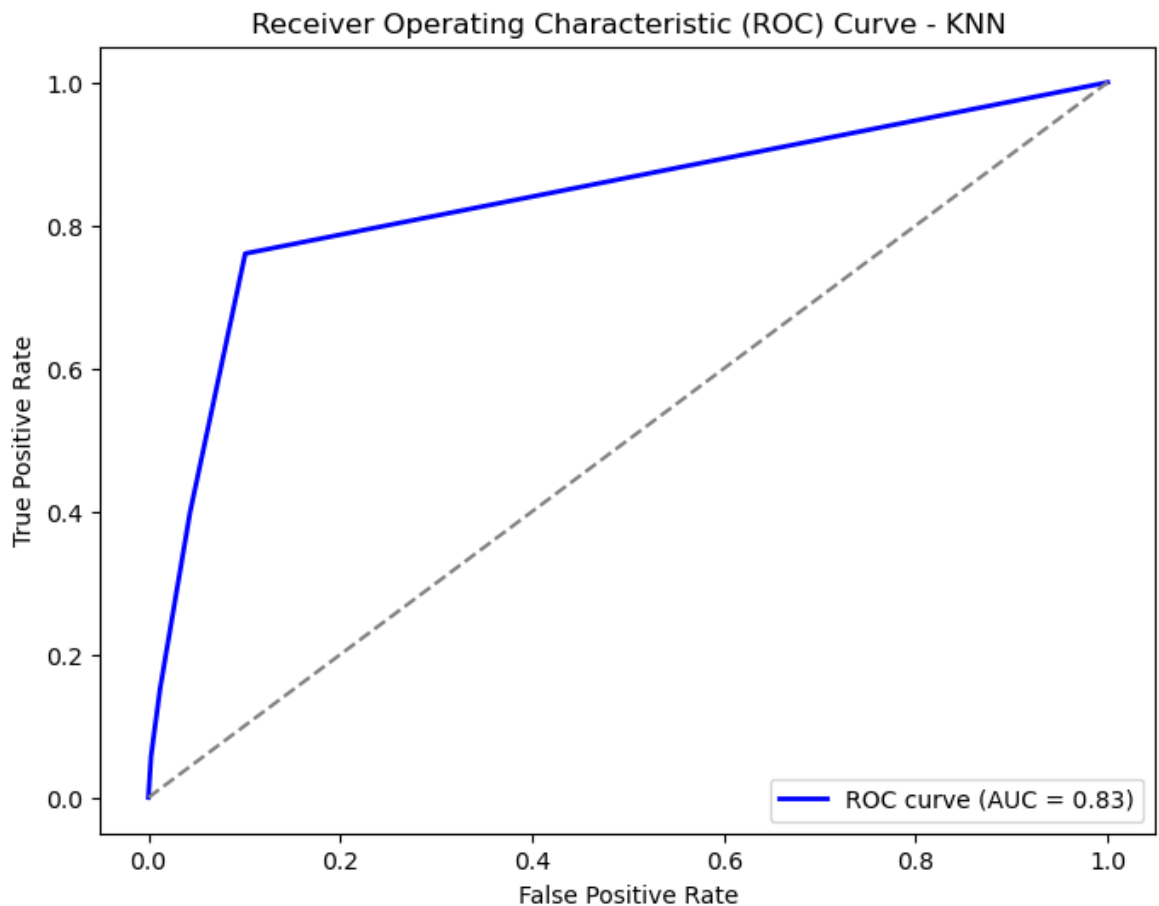
Precision: 0.94

Recall: 0.95

F-measure: 0.95

Specificity: 0.99





Model 5: SVM with Linear Kernel

Since linear SVM is $O(n^3)$, and the training dataset is large, it would take a lot of operations to train the model. Therefore we will only use 10000 datapoints total.

```
In [ ]: svm_train = training.sample(10000)
svm_X_train = svm_train[song_features]
svm_y_train = svm_train['popularity']
svm_X_test = song_data.drop(svm_train.index)[song_features]
svm_X_train, X_valid_LSVC, svm_y_train, svm_y_valid = train_test_split(svm_X_train,
```

```
svm_model = DecisionTreeClassifier()
svm_model.fit(svm_X_train, svm_y_train)
svm_y_pred = svm_model.predict(X_valid_LSVC)
svm_accuracy = accuracy_score(svm_y_valid, svm_y_pred)
print("Accuracy: " + str(svm_accuracy))

svm_auc = roc_auc_score(svm_y_valid, svm_y_pred)
print("AUC: " + str(svm_auc))
```

Accuracy: 0.9145
AUC: 0.5438740378895759

```
In [ ]: accuracy = accuracy_score(svm_y_valid, svm_y_pred)
conf_matrix = confusion_matrix(svm_y_valid, svm_y_pred)
precision = precision_score(svm_y_valid, svm_y_pred, average='weighted')
recall = recall_score(svm_y_valid, svm_y_pred, average='weighted')
f1 = f1_score(svm_y_valid, svm_y_pred, average='weighted')

tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
```



```

print("Confusion Matrix:")
print(conf_matrix)
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F-measure: {f1:.2f}")
print(f"Specificity: {specificity:.2f}")

plt.figure(figsize=(4, 3))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted Low', 'Predicted High'],
            yticklabels=['Actual Low', 'Actual High'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - SVM')
plt.show()

svm_y_pred_proba = svm_model.predict_proba(X_valid_LSVC)[: , 1]

fpr, tpr, thresholds = roc_curve(svm_y_valid, svm_y_pred_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - SVM')
plt.legend(loc='lower right')
plt.show()

```

Confusion Matrix:

```
[[1817  96]
 [ 75  12]]
```


Accuracy: 0.91


Precision: 0.92

Recall: 0.91

F-measure: 0.92

Specificity: 0.95





Model 6: Feedforward Neural Network

```
In [ ]: nn_model = Sequential()

nn_model.add(Dense(64, activation='relu', input_shape=(X_train.shape[1],)))
nn_model.add(Dense(32, activation='relu'))
nn_model.add(Dense(1, activation='sigmoid'))

nn_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

nn_model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_valid, y_valid))

eval_metrics = nn_model.evaluate(X_valid, y_valid)
print(f"Accuracy: {eval_metrics[1]}")

nn_y_pred = nn_model.predict(X_valid)

nn_auc = roc_auc_score(y_valid, nn_y_pred)
print(f"AUC: {nn_auc}")
```

```

Epoch 1/10
4655/4655 [=====] - 12s 2ms/step - loss: 0.1495 - accuracy: 0.9611 - val_loss: 0.1495 - val_accuracy: 0.9601
Epoch 2/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1453 - accuracy: 0.9615 - val_loss: 0.1489 - val_accuracy: 0.9601
Epoch 3/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1441 - accuracy: 0.9615 - val_loss: 0.1488 - val_accuracy: 0.9601
Epoch 4/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1436 - accuracy: 0.9616 - val_loss: 0.1479 - val_accuracy: 0.9601
Epoch 5/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1428 - accuracy: 0.9616 - val_loss: 0.1488 - val_accuracy: 0.9601
Epoch 6/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1422 - accuracy: 0.9616 - val_loss: 0.1485 - val_accuracy: 0.9601
Epoch 7/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1416 - accuracy: 0.9616 - val_loss: 0.1500 - val_accuracy: 0.9601
Epoch 8/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1409 - accuracy: 0.9616 - val_loss: 0.1469 - val_accuracy: 0.9600
Epoch 9/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1404 - accuracy: 0.9616 - val_loss: 0.1466 - val_accuracy: 0.9601
Epoch 10/10
4655/4655 [=====] - 11s 2ms/step - loss: 0.1398 - accuracy: 0.9616 - val_loss: 0.1459 - val_accuracy: 0.9600
1164/1164 [=====] - 2s 2ms/step - loss: 0.1459 - accuracy: 0.9600
Accuracy: 0.9600118398666382
1164/1164 [=====] - 2s 2ms/step
AUC: 0.7837994838216982

```

```

In [ ]: nn_y_pred_proba = nn_model.predict(X_valid)

fpr, tpr, thresholds = roc_curve(y_valid, nn_y_pred_proba)
nn_auc = auc(fpr, tpr)


plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % nn_auc)
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - Neural Network')
plt.legend(loc='lower right')
plt.show()

```

```

1/1164 [.....] - ETA: 44s1164/1164 [=====]
=====] - 1s 1ms/step

```



Model Performance

```
In [ ]: model_performance_accuracy = pd.DataFrame({
    'Model': ['LogisticRegression', 'RandomForestClassifier', 'KNeighborsClassifier'],
    'Accuracy': [lr_accuracy, rf_accuracy, knn_accuracy, dt_accuracy, svm_accuracy]
})

model_performance_AUC = pd.DataFrame({
    'Model': ['LogisticRegression', 'RandomForestClassifier', 'KNeighborsClassifier'],
    'AUC': [lr_auc, rf_auc, knn_auc, dt_auc, svm_auc]
})
```

```
In [ ]: model_performance_accuracy.sort_values(by = "Accuracy", ascending = False)
```

```
Out[ ]:
```

| | Model | Accuracy |
|---|------------------------|----------|
| 1 | RandomForestClassifier | 0.986196 |
| 3 | DecisionTreeClassifier | 0.967182 |
| 0 | LogisticRegression | 0.960066 |
| 2 | KNeighborsClassifier | 0.954506 |
| 4 | LinearSVC | 0.914500 |

```
In [ ]: model_performance_AUC.sort_values(by = "AUC", ascending = False)
```

Out[]:

| | Model | AUC |
|---|------------------------|----------|
| 3 | DecisionTreeClassifier | 0.843047 |
| 1 | RandomForestClassifier | 0.842960 |
| 2 | KNeighborsClassifier | 0.569936 |
| 4 | LinearSVC | 0.543874 |
| 0 | LogisticRegression | 0.500000 |

Results - We looked at data from songs on Spotify. We used details like the music's key, mood, and dance style to guess how popular a song might be (if it's rated higher than 70 out of 100) and didn't rely on things like the artist's name, music type, or when the song was released. - The best method we used was the Random Forest Classifier. It guessed accurately about 98 out of 100 times and scored about 84 out of 100 when measuring its performance. The next best method was the Decision Tree Classifier. It got it right about 96 out of 100 times and scored around 84 out of 100 in its performance.

References: 1. Halilovic, I. (2021, July 30). Markdown for Jupyter notebooks cheatsheet - Inge Halilovic - Medium. Medium. <https://ingeh.medium.com/markdown-for-jupyter-notebooks-cheatsheet-386c05aeebed> 2. Scribbr. (2021, July 30). Free APA citation Generator | with Chrome Extension - Scribbr. <https://www.scribbr.com/citation/generator/apa/> 3. Zach. (2023). How to create a distribution plot in Matplotlib. Statology. <https://www.statology.org/matplotlib-distribution-plot/> 4. seaborn.displot — seaborn 0.12.2 documentation. (n.d.). <https://seaborn.pydata.org/generated/seaborn.displot.html> 5. Lau, C. H. (2021, December 7). 5 steps of a Data Science Project Lifecycle - towards Data Science. Medium. <https://towardsdatascience.com/5-steps-of-a-data-science-project-lifecycle-26c50372b492> 6. Are hit songs becoming less musically diverse? (n.d.). The Pudding. <https://pudding.cool/2018/05/similarity/> 7. Nasreldin, M. (2018, July 2). Song Popularity Predictor - towards Data science. Medium. <https://towardsdatascience.com/song-popularity-predictor-1ef69735e380> 8. Seaborn.Heatmap — seaborn 0.13.0 documentation. (n.d.). Pydata.org. Retrieved October 3, 2023. <https://seaborn.pydata.org/generated/seaborn.heatmap.html> 9. IBM documentation. (2021, April 8). Ibm.com. <https://www.ibm.com/docs/en/watson-studio-local/1.2.3?topic=notebooksmarkdown-jupyter-cheatsheet> 10. Spotify Tracks DB. (2019, July 23). Kaggle. <https://www.kaggle.com/zaheenhamidani/ultimate-spotify-tracks-db#SpotifyFeatures.csv>